

# Elm

IN ACTION

Richard Feldman



MEAP



MANNING



**MEAP Edition  
Manning Early Access Program  
Elm in Action  
Version 2**

Copyright 2016 Manning Publications

For more information on this and other Manning titles go to  
[www.manning.com](http://www.manning.com)

# welcome

---

Thanks for buying the MEAP of *Elm in Action*! I hope you find this a wonderfully practical journey into a delightful language.

As an early adopter of Elm, I felt like a wide-eyed sprinter barreling down an exciting path that had not yet been paved. Now that I've had time to catch my breath, I'm eager to help pave that road for future travelers. My goal for this book is for you to enjoy learning Elm even more than I did...and I loved it so much I'm writing a book about it!

More than anything, I hope this book gives you the confidence to use Elm professionally. As much as I enjoyed using Elm on a side project, it's brought me even more joy at work. When I think back on what motivated my team to try it, it was the practical benefits: reliability, a quicker development loop, and a lower maintenance burden.

The most authentic way I know to convey the benefits of building an application in Elm is...well, to build an application in Elm. That's exactly what we'll be doing over the course of this book: developing an application from start to finish, learning the language along the way. By the end you'll have built, refactored, and tested an Elm code base. If using it at work sounds appealing, you'll be able to convey its benefits based on firsthand experience!

Throughout the MEAP process I intend both to release new chapters as I finish them, and to revise past chapters based on feedback from readers like you. I'd love it if you could leave some comments in the Author Online Forum. All feedback is helpful! Mentioning which parts felt difficult tells me where to focus my revision time, and mentioning your favorite parts tells me what not to change.

Thanks again, and enjoy!

—Richard Feldman

# *brief contents*

---

## **PART 1: GETTING STARTED**

- 1 Welcome to Elm*
- 2 Your First Elm Application*
- 3 Compiler as Assistant*

## **PART 2: PRODUCTION-GRADE ELM**

- 4 Talking to Servers*
- 5 Talking to JavaScript*
- 6 Testing*
- 7 Tools*

## **PART 3: BUILDING BIGGER**

- 8 Single-Page Applications*
- 9 Scaling Elm Code*
- 10 Performance Optimization*

*Appendix A: Getting Set Up*

## 1

# Welcome to Elm

## This chapter covers

- How and why to introduce Elm to a project
- Using `elm-repl`
- Building expressions
- Writing and importing functions
- Working with collections

Back in 2014 I set out to rewrite a side project, and ended up with a new favorite programming language. Not only was the rewritten code faster, more reliable, and easier to refactor, writing it was the most fun I'd had in over a decade writing code. Ever since that project, I've been hooked on Elm.

The rewrite in question was a writing application I'd built even longer ago, in 2011. Having tried out several writing apps over the course of writing a novel, and being satisfied with none, I decided to scratch my own itch and build my dream writing app. I called it Dreamwriter.

For those keeping score: yes, I was indeed writing code in order to write prose better.

Things went well at first. I built the basic Web app, started using it, and iterated on the design. Months later I'd written over fifty thousand words in Dreamwriter. If I'd been satisfied with that early design, the story might have ended there. However, users always want a better experience...and when the user and the developer are the same person, further iteration is inevitable.

The more I revised Dreamwriter, the more difficult it got to maintain. I'd spend hours trying to reproduce bugs that knocked me out of my writing groove. At some point the copy and paste functions stopped working, and I found myself resorting to the browser's developer tools whenever I needed to move paragraphs around.

Right around when I'd decided to scrap my unreliable code base and do a full rewrite, a blog post crossed my radar. After reading it I knew three things:

1. The Elm programming language compiled to JavaScript, just like Babel or CoffeeScript. (I already had a compile step in my build script, so this was familiar territory.)
2. Elm used the same rendering approach as React.js—which I had recently grown to love—except Elm had rendering benchmarks that outperformed React's!
3. Elm's compiler would catch a lot of the errors I'd been seeing before they could harm me in production. I did not yet know just how many it would catch.

I'd never built anything with a functional programming language like Elm before, but I decided to take the plunge. I didn't really know what I was doing, but the compiler's error messages kept picking me up whenever I stumbled. Eventually I got the revised version up and running, and began to refactor.

The refactoring experience blew me away. I revised the Elm-powered Dreamwriter gleefully, even recklessly—and no matter how dramatic my changes, the compiler always had my back. It would point out whatever corner cases I'd missed, and I'd go through and fix them. As soon as the code compiled, lo and behold, everything worked again. I felt *invincible*.

I related my Elm experience to my coworkers at NoRedInk, and they were curious but understandably cautious. How could we find out if the team liked it without taking a big risk? A full rewrite may have been fine for Dreamwriter, but it would have been crazy to attempt that for our company's entire front-end.

So we introduced Elm gently, by rewriting just one portion of one production feature in Elm. It went well, so we did a bit more. And then more.

Today our front-end is as Elm-powered as we can make it, and our team has never been happier. Our test suites are smaller, yet our product is more reliable. Our feature set has grown more complex, yet refactoring remains delightful. We swap stories with other companies using Elm about how long our production code has run without throwing a single runtime exception.

In this book we'll explore all of these benefits.

After learning some basics, we'll build an Elm Web application the way teams typically do: ship a basic version that works, but which has missing features and some technical debt. As we advance through the chapters, we'll expand and refactor this application, adding features and paying off technical debt as we learn more about Elm. We'll debug our business logic. We'll even interoperate with JavaScript, using its vast library ecosystem to avoid reinventing the wheel. By the end of the book we will have transformed our application into a more featureful product, with a more maintainable code base, than the one we initially shipped.

With any luck, we'll have a lot of fun doing it.

Welcome to Elm!

## 1.1 How Elm Fits In

Elm can be used either as a replacement for in-browser JavaScript code, or as a complement to it. You write some .elm files, run them through Elm’s compiler, and end up with plain old .js files that the browser runs as normal. If you have separate stylesheets that you use alongside JavaScript, they’ll work the same way alongside Elm.




Figure 1.1 Elm files are compiled to plain old JavaScript files

The appropriate Elm-to-JavaScript ratio can vary by project. Some projects may want primarily JavaScript and only a touch of Elm for business logic or rendering. Others may want a great deal of Elm but just a pinch of JavaScript to leverage its larger ecosystem. No single answer applies to every project.

What distinguishes Elm from JavaScript is **maintainability**.

Handwritten JavaScript code is notoriously prone to runtime crashes like “undefined is not a function.” In contrast, Elm code has a reputation for never throwing runtime exceptions in practice. This is because Elm is built around a small set of simple primitives like expressions, immutable values, and managed effects. That design lets the compiler identify lurking problems just by scanning your source code. It reports these problems with such clarity that it has earned a reputation for user-friendliness even among programming legends.

*That should be an inspiration for every error message.*

**—John Carmack, after seeing one of Elm’s compiler errors**

Having this level of compiler assistance makes Elm code dramatically easier to refactor and debug, especially as code bases grow larger. There is an up-front cost to learning and adopting Elm, but you reap more and more maintainability benefits the longer the project remains in active development.

**TIP** Most teams that use Elm in production say they used a “planting the seed” approach. Instead of waiting for a big project where they could build everything in Elm from the ground up, they rewrote a small part of their existing JavaScript code base in Elm. This was low-risk and could be rolled back if things did not go as planned,

but having that small seed planted in production meant they could grow their Elm code at a comfortable pace from then on.

Although Elm is in many ways a simpler language than JavaScript, it's also much younger. The ecosystem is smaller, which means you'll encounter fewer problems where someone has already published an off-the-shelf solution. Elm code can interoperate with JavaScript code to piggyback that ecosystem, but Elm's design differs enough from JavaScript's that incorporating JavaScript libraries takes effort.

Balancing these tradeoffs depends on the specifics of a given project. Let's say you're on a team where people are comfortable with JavaScript but are new to Elm. Here are some projects I'd expect would benefit from learning and using Elm:

- Feature-rich Web applications whose code bases are large or will grow large
- Individual features that will be revised and maintained over an extended period of time
- Projects where most functionality comes from in-house code, not off-the-shelf libraries

In contrast, I'd stick to a more familiar language and toolset for projects like these:

- Proof-of-concept prototypes that will not be maintained long-term
- Time-crunched projects where learning a language is unrealistic given the deadline
- Projects that will consist primarily of gluing together off-the-shelf components

We'll explore these tradeoffs in more detail throughout the course of the book.

## 1.2 Expressions

To get our feet wet with Elm, let's tap into one of the most universal traits across the animal kingdom: the innate desire to *play*.

Researchers have developed many theories as to why we play, including to learn, to practice, to experiment, and of course for the pure fun of it. These researchers could get some high-quality data by observing a member of the *homo sapiens programmerus* species in its natural environment for play—the Read-Eval-Print Loop, or REPL.

You'll be using Elm's REPL to play as you take your first steps as an Elm programmer.

### 1.2.1 Using elm-repl

The Elm Platform includes a nice REPL called `elm-repl`, so if you have not installed the Elm Platform yet, head over to Appendix A to get hooked up.

Once you're ready, enter `elm-repl` at the terminal. You should see this prompt:

```
---- elm-repl 0.17.1 -----
:help for help, :exit to exit, more at https://github.com/elm-lang/elm-repl
-----
>
```



Alexander Graham Bell invented the telephone over a century ago. There was no customary greeting back then, so Bell suggested one: lift the receiver and bellow out a rousing “Ahoy!” Thomas Edison later proposed the alternative “Hello,” which stuck, and today programmers everywhere append “World” as the customary way to greet a new programming language.

Let’s spice things up a bit, shall we? Enter this at the prompt.

```
> "Ahoy, World!"
```

You should see this response from `elm-repl`:

```
"Ahoy, World!" : String
```

Congratulations, you are now an Elm programmer!

**NOTE** To focus on the basics, for the rest of this chapter we’ll omit the type annotations that `elm-repl` prints. For example, the previous code snippet would have omitted the `: String` portion of `"Ahoy, World!" : String`. We’ll get into these annotations in Chapter 3.

If you’re the curious sort, by all means feel free to play as we continue. Enter things that occur to you, and see what happens! Whenever you encounter an error you don’t understand yet, picture yourself as a tiger cub building intuition for physics through experimentation: adorable for now, but powerful in time.

## 1.2.2 Building Expressions

Let’s rebuild our `"Ahoy, World!"` greeting from two parts, and then play around from there. Try entering these into `elm-repl`.

### Listing 1.1 Combining Strings

```
> "Ahoy, World!"
"Ahoy, World!"

> "Ahoy, " ++ "World!"
"Ahoy, World!"

> "Pi is " ++ toString pi ++ " (give or take)" ❶
"Pi is 3.141592653589793 (give or take)"
```

❶ `toString` is a standalone function, not a method

In Elm, we use the `++` operator to combine strings, instead of the `+` operator JavaScript uses. At this point you may be wondering: Does Elm even have a `+` operator? What about the other arithmetic operators?

Let’s find out by experimenting in `elm-repl`!

### Listing 1.2 Arithmetic Expressions

```
> 1234 + 103
1337

> 12345 - (5191 * -15) ❶
90210

> 2 ^ 11
2048

> 49 / 10
4.9

> 49 // 10 ❷
4

> -5 % 2 ❸
1
```

- ❶ Nest expressions via parentheses
- ❷ Integer division (decimals get truncated)
- ❸ Remainder after integer division

Sure enough, Elm has both a ++ and a + operator. They are used for different things:

- The ++ operator is for appending. Using it on a number is an error.
- The + operator is for addition. It can *only* be used on numbers.

You will see this preference for **being explicit** often in Elm. If two operations are sufficiently different—in this case, adding and appending—Elm implements them separately, so each implementation can **do one thing well**.

## STRINGS AND CHARACTERS

Elm also distinguishes between strings and the individual UTF-8 *characters* that comprise them. Double quotes in Elm represent string literals, just like in JavaScript, but single quotes in Elm represent character literals.

Table 1.1 Strings and Characters

Elm Literal	Result
"a"	a string with a length of 1
'a'	a single character
"abc"	a string with a length of 3
'abc'	<b>error:</b> character literals must contain exactly 1 character
""	an empty string
' '	<b>error:</b> character literals must contain exactly 1 character

## COMMENTS

Comments also work a bit differently in Elm:

- Use `--` instead of `//` for inline comments
- Use `{- -}` instead of `/* */` for block comments

Let's see these in action!

### Listing 1.3 Characters, Comments, and Constants

```
> 'a' -- This is a comment. It will be ignored. ❶
'a'

> "a" {- This is a block comment. It will also be ignored. -} ❷
"a"

> milesPerHour = 88 ❸
88

> milesPerHour
88
```

- ❶ JavaScript comment: `//`
- ❷ JavaScript comment: `/* ... */`
- ❸ JavaScript: `const milesPerHour = 88;`

## NAMING VALUES WITH CONSTANTS

In the last two lines of code above, we did something new: we assigned the *constant* `milesPerHour` to the value `88`.

**DEFINITION** A *constant* assigns a name to a value. Once assigned, this name cannot be later reassigned to a different value in the same scope.

There are a few things to keep in mind when naming constants.

- The name must begin with a lowercase letter. After that it can be a mix of letters, numbers, underscores, and apostrophes.
- By convention, all letters should be in one uninterrupted sequence. For example, `map4` is a reasonable name, but `map4ever` is not, because the sequence of letters is interrupted by the `4`.
- Because of the previous two rules, you should never use `snake_case` or `SCREAMING_SNAKE_CASE` to name constants. Use `camelCase` instead.
- It can be tempting to name a new constant by recycling an old name and slapping an apostrophe or underscore on the end, like defining `list` and then defining `list'` or `list_` shortly after. This is a code smell which leads to similar readability problems as un-descriptive single-letter identifiers do in JavaScript. Take the time to pick a more

descriptive name; whoever reads your code next will be glad you did!

- If you absolutely must know whether the compiler will accept `richard's_rad_THING` as a valid constant name, remember: what happens in `elm-repl` stays in `elm-repl`.

## ASSIGNING CONSTANTS TO EXPRESSIONS

Not only can you assign constants to literal values, you can also assign them to *expressions*.

**DEFINITION** An *expression* is anything that evaluates to a single value.

Here are some expressions we've seen so far.

Expression	Evaluates to
"Ahoy, " ++ "World!"	"Ahoy, World!"
2 ^ 11	2048
pi	3.141592653589793
42	42

**NOTE** Since an expression is anything that evaluates to a value, literal values like "Ahoy, World!" and 42 are expressions too—just expressions that have already been fully evaluated.

Expressions are the basic building block of Elm applications. This is different from JavaScript, which offers many features as *statements* instead of expressions.

Consider these two lines of JavaScript code.

```
label = (num > 0) ? "positive" : "negative" // ternary expression
label = if (num > 0) { "positive" } else { "negative" } // if-statement
```

The first line is ternary *expression*. Being an expression, it evaluates to a value, and JavaScript happily assigns that value to `label`.

The second line is an *if-statement*, and since statements do not evaluate to values, trying to assign it to `label` yields a syntax error.

This distinction does not exist in Elm, as Elm programs express logic using expressions only. As such, Elm has *if-expressions* instead of *if-statements*. As we will see in Chapter 2, every Elm application is essentially one big expression built up from lots of smaller ones!

### 1.2.3 Booleans and Conditionals

There aren't terribly many booleans out there—just the two, really—and working with them in Elm is similar to working with them in JavaScript. There are a few differences, though.

- You write `True` and `False` instead of `true` and `false`
- You write `/=` instead of `!=`
- To negate values, you use Elm's `not` function instead of JavaScript's `!` prefix

Let's try them out!

#### Listing 1.4 Boolean Expressions

```
> pi == pi ❶
True      ❷

> pi /= pi ❸
False     ❹

> not (pi == pi) ❺
False

> pi <= 0 || pi >= 10
False

> 3 < pi && pi < 4 ❻
True
```

- ❶ JavaScript: `pi === pi`
- ❷ JavaScript: `true`
- ❸ JavaScript: `pi !== pi`
- ❹ JavaScript: `false`
- ❺ JavaScript: `!(pi === pi)`
- ❻ `3 < pi < 4` would be an error

Now let's say it's a lovely afternoon at the North Pole, and we're in Santa's workshop writing a bit of UI logic to display how many elves are currently on vacation. The quick-and-dirty approach would be to add the string " elves" after the number of vacationing elves, but then when the count is 1 we'd display "1 elves", and we're better than that.

Let's polish our user experience with the *if-expression* shown in Figure 1.2.




Figure 1.2 Comparing an Elm *if-expression* to a JavaScript *ternary*

Like JavaScript ternaries, Elm *if-expressions* require three ingredients:

1. A condition
2. A branch to evaluate if the condition passes
3. A branch to evaluate otherwise

Each of these ingredients must be expressions, and the *if-expression* itself evaluates to the result of whichever branch got evaluated. You'll get an error if any of these three ingredients are missing, so make sure to specify an `else` branch every time!

**NOTE** There is no such thing as “truthiness” in Elm. Conditions can be either `True` or `False`, and that's it. Life is simpler this way.

Now let's say we modified our pluralization conditional to include a third case:

- If we have one Elf, evaluate to `"elf"`
- Otherwise, if we have a positive number of elves, evaluate to `"elves"`
- Otherwise, we must have a negative number of elves, so evaluate to `"anti-elves"`

In JavaScript you may have used `else if` to continue branching conditionals like this. It's common to use `else if` for the same purpose in Elm, but it's worth noting that `else if` in Elm is nothing more than a stylish way to combine the concepts we learned a moment ago.

Check it out!

#### Listing 1.6 Using `else if`

```

if elfCount == 1 then           ❶
  "elf"                         ❶
else                             ❶
  (if elfCount >= 0 then "elves" else "anti-elves") ❶

if elfCount == 1 then           ❷
  "elf"                         ❷
else (if elfCount >= 0 then      ❷
  "elves"                       ❷
else                             ❷
  "anti-elves")                 ❷

if elfCount == 1 then           ❸
  "elf"                         ❸
else if elfCount >= 0 then      ❸
  "elves"                       ❸
else                             ❸
  "anti-elves"                 ❸

```

- ❶ Use an *if-expression* inside `else`
- ❷ Rearrange some whitespace
- ❸ Drop the parentheses

This works because the `else` branch of an *if-expression* must be an expression, and it just so happens that *if-expressions* themselves are expressions. As shown in Figure 1.3, all it takes is putting an *if-expression* after another one's `else`, and *voilà!* Additional branching achieved.

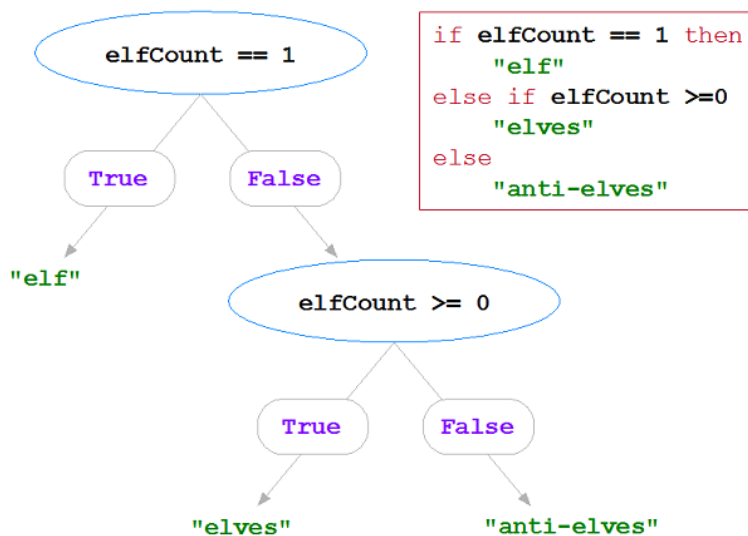


Figure 1.3 The *else if* technique: use an *if-expression* as the *else branch* of another *if-expression*

Nesting expressions for fun and profit is a recurring theme in Elm, and we'll see plenty more recipes like `else if` throughout the book.

Chapter 3 will add a powerful new conditional to our expression toolbox, one which has no analogue in JavaScript: the *case-expression*.

## 1.3 Functions

Earlier we wrote this expression:

```
elfLabel = if vacationingElves == 1 then "elf" else "elves"
```

Suppose it turns out a general-purpose singular/plural labeler would be really useful, and we want to reuse similar logic across the code base at Santa's Workshop. Search results might want to display "1 result" and "2 results" as appropriate, for example.

We can write a *function* to make this pluralization logic easily reusable.

**DEFINITION** Elm *functions* represent reusable logic. They are not objects. They have no fields, no prototypes, and no ability to store state. All they do is accept values as arguments, and then return a value.

If you thought expressions would be a recurring theme in Elm, wait 'til you see functions!

### 1.3.1 Defining Functions

Let's define our first function: `isOdd`. It will take a number and then:

- Return `True` if the number is odd

- Return `False` if the number is even

We can define `isOdd` in `elm-repl` and try it out right away.

### Listing 1.7 Defining a function

```
> isOdd num = num % 2 == 1 ❶
<function>

> isOdd 5 ❷
True

> (isOdd 5) ❸
True

> isOdd (5 + 1) ❹
False
```

- ❶ JavaScript: `function isOdd(num) { return num % 2 === 1; }`
- ❷ JavaScript: `isOdd(5)`
- ❸ JavaScript: `(isOdd(5))`
- ❹ JavaScript: `isOdd(5 + 1)`

As you can see, in Elm we put the function parameter name before the `=` sign. We also don't surround the function body with `{ }`. And did you notice the `return` keyword is nowhere to be seen? That's because Elm doesn't have one! In Elm, a function body is a single expression, and since an expression evaluates to a single value, Elm uses that value as the function's return value. This means all Elm functions return values!

For our `isOdd` function, the expression `num % 2 == 1` serves as the function's body, and provides its return value.

### Refactoring out an early return

In JavaScript, `return` is often used to exit a function early. This is harmless when used responsibly, but can lead to unpleasant surprises when used in the middle of large functions. Elm does not support these unpleasant surprises, because it has no `return` keyword.

Let's refactor the early `return` out of this function:

```
function capitalize(str) {
  if (!str) {
    return str; ❶
  }

  return str[0].toUpperCase() + str.slice(1);
}
```

- ❶ Early return

Without making any other changes, we can refactor this early `return` into a ternary:



```
function capitalize(str) {
  return !str ? str : str[0].toUpperCase() + str.slice(1);
}
```

*Poof!* There it goes. Since JavaScript's ternaries are semantically analogous to Elm's *if-expressions*, this code is now much more straightforward to rewrite in Elm. More convoluted JavaScript functions may require more steps than this, but it is always possible to untangle them into plain old conditionals.

Removing an early `return` is one of many quick refactors you can do to ease the transition from legacy JavaScript to Elm, and we'll look at more of them throughout the book. When doing these, do not worry if the intermediate JavaScript code looks ugly! It's intended to be a stepping stone to nicer Elm code, not something to be maintained long-term.

Let's use what we just learned to generalize our previous `elf-labeling` expression into a reusable `pluralize` function. Our function this time will have a longer definition than last time, so let's use multiple lines to give it some breathing room. In `elm-repl`, you can enter multiple lines by adding `\` to the end of the first line and indenting the next line.

**NOTE** Indent with spaces only! Tab characters are syntax errors in Elm.

#### Listing 1.8 Using multiple REPL lines

```
> pluralize singular plural count = \
|   if count == 1 then singular else plural ❶
<function>

> pluralize "elf" "elves" 3 ❷
"elves"

> pluralize "elf" "elves" (round 0.9) ❸
"elf"
```

- ❶ Don't forget to indent!
- ❷ No commas between arguments!
- ❸ `(round 0.9)` returns 1

When passing multiple arguments to an Elm function, separate the arguments with whitespace and not commas. That last line of code is an example of passing the result of one function call, namely `round 0.9`, as an argument to another function. (Think about what would happen if we did not put parentheses around `(round 0.9)`...how many arguments would we then be passing to `pluralize`?)

### 1.3.2 Importing Functions

So far we've only used basic operators and functions we wrote ourselves. Now let's expand our repertoire of functions by using one from an external module.

**DEFINITION** A *module* is a named collection of Elm functions and other values.

The `String` module is one of the core modules that ships with Elm. Additional modules can be obtained from Elm's official package repository, copy-pasting code from elsewhere, or through a back-alley rendezvous with a shadowy figure known as Dr. Deciduous. Chapter 4 will cover how to do the former, but neither the author nor Manning Publications endorses obtaining Elm modules through a shadowy back-alley rendezvous.

Let's import the `String` module and try out two of its functions, `toLowerCase` and `toUpperCase`.

### Listing 1.9 Importing functions

```
> import String
> String.toLowerCase "Why don't you make TEN louder?"
"why don't you make ten louder?"

> String.toUpperCase "These go to eleven."
"THESE GO TO ELEVEN."
```

Observant readers may note a striking resemblance between Elm's `String.toUpperCase` function and the `toUpperCase()` method one finds on JavaScript strings. This is the first example of a pattern we will encounter many times!

JavaScript has several ways of organizing string-related functionality: fields on a string, methods on a string, or methods on the `String` global itself.

In contrast, Elm strings have neither fields nor methods. The `String` module houses the standard set of string-related features, and exposes them in the form of plain old functions like `toLowerCase` and `toUpperCase`.

**Table 1.2 String Functionality Comparison**

JavaScript	Elm
"storm".length	<code>String.length</code> "storm"
"dredge".toUpperCase()	<code>String.toUpperCase</code> "dredge"
<code>String.fromCharCode</code> (something)	<code>String.fromCharCode</code> something

Not only is this organizational pattern consistent within the `String` module, it's consistent across Elm. Want a standard date-related feature? Look no further than the functions in the `Date` module. Regular Expression functions? Hit up the `Regex` module.

Methods are never the answer in Elm; over here it's all vanilla functions, all the time.

**TIP** Complete documentation for `String`, `Date`, `Regex`, and other tasty modules can be found in the *core* section of the [package.elm-lang.org](https://package.elm-lang.org) website.

We'll learn more about modules in the coming chapters, including how to write our own!

## USING STRING.FILTER TO FILTER OUT CHARACTERS

Another useful function in the `String` module is `filter`. It lets us filter out unwanted characters from a string, such as non-numeric digits from a phone number.

To do this, we must give `filter` a function which specifies which characters to keep. The function will take a single character as an argument and return `True` if we should keep that character or `False` if we should chuck it. Figure 1.4 illustrates using `String.filter` to remove dashes from a US telephone number.

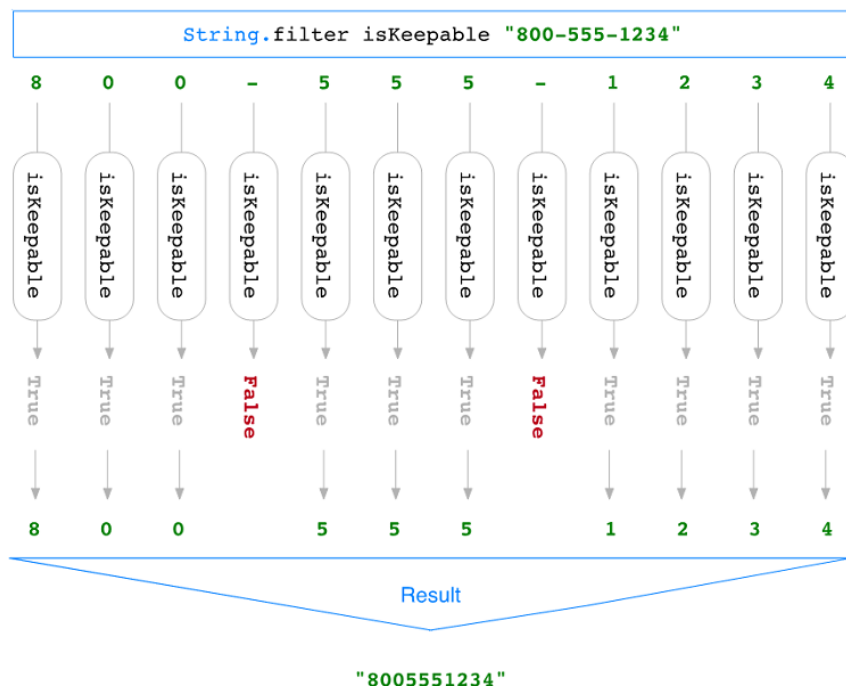


Figure 1.4 Using `String.filter` to remove dashes from a US phone number

As in JavaScript, Elm functions are first-class values that can be passed around just like any other value. This lets us provide `filter` with the function it expects by defining that function and then passing it in as a plain old argument.

### Listing 1.10 Filtering with a named function

```
> import String
> isKeepable character = character /= '-' ❶
<function>

> isKeepable 'z'
True
```

```
> isKeepable '-'
False

> String.filter isKeepable "800-555-1234" ❷
"8005551234"
```

- ❶ A function describing which characters to keep
- ❷ Passing our function to `String.filter`

This normalizes telephone numbers splendidly. Alexander Graham Bell would be proud!

`String.filter` is one of the *higher-order functions* (that is, functions which accept other functions as arguments) that Elm uses to implement customizable logic like this.

### 1.3.3 Creating scope with *let*-expressions

Let's say we find ourselves removing dashes from phone numbers so often, we want to make a reusable function for it. We can do that with our trusty `isKeepable` function:

```
withoutDashes str = String.filter isKeepable str
```

This works, but in a larger Elm program, it might be annoying having `isKeepable` in the global scope like this. After all, its implementation is only useful to `withoutDashes`. Can we avoid globally reserving such a nicely self-documenting name?

Absolutely! We can scope `isKeepable` to the implementation of `withoutDashes` using a *let-expression*.

**DEFINITION** A *let-expression* adds locally scoped constants to an expression.

Figure 1.5 shows how we can implement `withoutDashes` using a single *let-expression*.

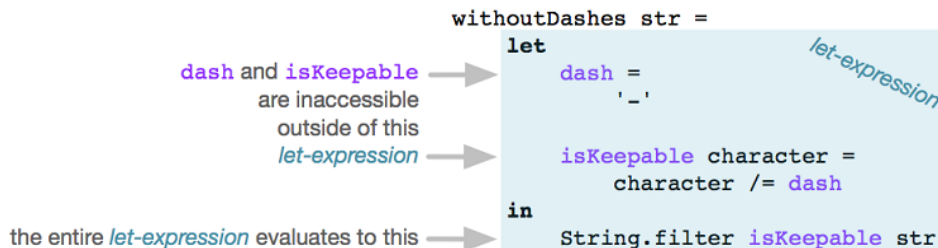


Figure 1.5 Anatomy of the wild *let-expression*

The above code does very nearly the same thing as entering the following in `elm-repl`:

```
> import String
> dash = '-'
> isKeepable character = character /= dash
> withoutDashes str = String.filter isKeepable str
```

In both versions, the implementation of `withoutDashes` boils down to `String.filter isKeepable str`. The only difference between the two is the scope of `dash` and `isKeepable`.

- In the `elm-repl` version above, `dash` and `isKeepable` are in the global scope.
- In Figure 1.5, `dash` and `isKeepable` are scoped locally to the *let-expression*.

Whenever you see a *let-expression*, you can mentally replace it with the part after its `in` keyword—in this case, `String.filter isKeepable str`. All the values between `let` and `in` are intermediate constants that are no longer in scope once the expression after `in` gets evaluated.

**NOTE** The indentation in Figure 1.5 is no accident! In a multiline *let-expression*, the `let` and `in` keywords must be at the same indentation level, and all other lines in the *let-expression* must be indented further than they are.

Anywhere you'd write a normal expression, you can swap in a *let-expression* instead. Because of this, you don't need to learn anything new to define locally-scoped constants inside function bodies, branches of *if-expressions*, or anywhere else.

Wherever you want some local scope, reach for a refreshing *let-expression*!

### 1.3.4 Anonymous Functions

Anonymous functions in Elm work the same way they do in JavaScript: like named functions but without the name.

#### Listing 1.11 Named and Anonymous Functions

```
function area(w, h) { return w * h; } ❶
function(w, h) { return w * h; }      ❷
area w h = w * h                      ❸
\w h -> w * h                         ❹
```

- ❶ JavaScript named function
- ❷ JavaScript anonymous function
- ❸ Elm named function
- ❹ Elm anonymous function

Elm's anonymous functions differ from its named functions in three ways.

1. They have no names
2. They begin with a `\`
3. Their parameters are followed by `->` instead of `=`

Once defined, anonymous functions and named functions work the same way; you can always use the one in place of the other. For example, the following do exactly the same thing:

```
isKeepable char = char /= '-'
isKeepable = \char -> char /= '-'
```

Let's use an anonymous function to call `String.filter` in one line instead of two, then see if we can improve the business logic! For example, we can try using `Char.isDigit` to cast a wider net, filtering out any non-digit characters instead of just dashes.

### Listing 1.12 Filtering with anonymous functions

```
> import String
> String.filter (\char -> char /= '-') "800-555-1234"
"8005551234"

> String.filter (\char -> char /= '-') "(800) 555-1234"
"(800) 5551234" ❶

> import Char
> String.filter (\char -> Char.isDigit char) "(800) 555-1234"
"8005551234" ❷

> String.filter Char.isDigit "(800) 555-1234" ❸
"8005551234"
```

- ❶ Our simple filter fell short here
- ❷ Much better!
- ❸ Refactor of previous approach

Anonymous functions are often used with higher-order functions like `String.filter`.

## 1.3.5 Operators

So far we've seen functions such as `String.filter`, as well as operators such as `++`, `-`, and `==`. How do operators and functions relate?

As it turns out, Elm's operators *are* functions! There are a few things that distinguish operators from normal functions:

- Operators must always accept exactly two arguments—no more, no fewer.
- Normal functions have names that begin with a letter. You typically call them by writing the name of the function followed by its arguments. This is *prefix-style* calling.
- Operators have names that contain neither letters nor numbers. You typically call them by writing the first argument, followed by the operator, followed by the second argument. This is *infix-style* calling.

Wrapping an operator in parentheses treats it as a normal function—*prefix-style* calling and all!




Figure 1.6 Calling the `(-)` operator in both *infix-style* and *prefix-style*

Let's play with some operators in `elm-repl`:

### Listing 1.13 Operators are functions

```
> (/)
<function>

> divideBy = (/)
<function>

> 7 / 2      ❶
3.5

> (/) 7 2    ❷
3.5

> divideBy 7 2
3.5

> 7 `divideBy` 2  ❸
3.5
```

- ❶ *infix-style* calling
- ❷ *prefix-style* calling
- ❸ backticks call normal functions in *infix-style*

As the final entry in Listing 1.13 demonstrates, you can also call normal functions in *infix-style* by surrounding them with backticks.

### OPERATOR PRECEDENCE

Try entering an expression involving both arithmetic operators and `(==)` into `elm-repl`:

```
> 3 + 4 == 8 - 1
True : Bool
```

Now consider how we'd rewrite this expression in *prefix-style*:

```
> (==) ((+) 3 4) ((-) 8 1)
True : Bool
```

Notice anything about the order in which these operators appear? Reading the *infix-style* expression from left to right, you see `+` first, then `==`, and finally `-`. In the *prefix-style* expression, the order is different: first you see `==`, then `+`, and finally `-`. Why is this?

They get reordered because `(==)`, `(+)`, and `(-)` have different *precedence* values.




Figure 1.7 `(==)` gets evaluated after `(+)` and `(-)` because it has lower precedence

**DEFINITION** In any expression containing multiple operators, the operators with higher *precedence* get evaluated before those with lower precedence. This only applies to *infix-style* calls, as all *prefix-style* calls implicitly have the same precedence.

There isn't much formal documentation on operators' relative precedence values, but operators that appear in many programming languages (such as the `(==)`, `(+)`, and `(-)` operators) tend to work similarly in Elm to how they do everywhere else.

We'll dig into operator precedence in greater depth in later chapters.

### NORMAL FUNCTION CALLS HAVE TOP PRECEDENCE

Here are two ways to write the same thing:

```
> negate 1 + negate 5
-6

> (negate 1) + (negate 5)
-6
```

These two are equivalent because normal function calls have higher precedence than any operator. This means any time you want to pass the results of two normal function calls to an operator, you won't need to add any parentheses! You'll still get the result you wanted.



## OPERATOR ASSOCIATIVITY

Besides precedence, the other factor that determines evaluation order for operators called in *infix-style* is whether the operators are *left-associative*, *right-associative*, or *non-associative*. Every operator is one of these.

An easy way to think about operator associativity is in terms of where the implied parentheses go. Infix expressions involving left-associative operators, such as arithmetic operators, have implied parentheses that cluster on the left:

**Table 1.3 Implied parentheses for the (-) operator**

Parentheses Shown	Expression	Result
none	<code>10 - 6 - 3</code>	1
assuming <i>left-associative</i>	<code>((10 - 6) - 3)</code>	1
assuming <i>right-associative</i>	<code>(10 - (6 - 3))</code>	7

If (-) were right-associative, `10 - 6 - 3` would have parentheses clustering on the right, meaning it would evaluate to `(10 - (6 - 3))` and the undesirable result of `10 - 6 - 3 == 7`. Good thing arithmetic operators are left-associative!

Non-associative operators cannot be chained together. For example, `foo == bar == baz` does not result in clustered parentheses, it results in an error!

In Chapter 7, we'll see how to define our own operators, including specifying their associativity and precedence values.

## 1.4 Collections

Elm's most basic collections are lists, records, and tuples. Each has varying degrees of similarity to JavaScript's arrays and objects, but one way in which they differ from JavaScript collections is in that Elm collections are always *immutable*.

**DEFINITION** An *immutable* value cannot be modified in any way once created.

This is in contrast to JavaScript, where some values (like strings and numbers) are immutable, but collections (like arrays and objects) can be mutated.

### 1.4.1 Lists

An Elm list has many similarities to a JavaScript array.

- You can create one with a square bracket literal, e.g. `[ "one fish", "two fish" ]`
- You can ask for its first element
- You can ask for its length
- You can iterate over its elements in various ways

An Elm list does have some differences, though.

- It is immutable
- It has no fields or methods. You work with it using functions from the `List` module.
- Because it is a *linked list*, you can ask for its first element, but not for other individual elements. (If you need to ask for elements at various different positions, you can first convert from an Elm `List` to an Elm `Array`. We'll discuss Elm Arrays in Chapter 3.)
- You can combine it with another list using the `++` operator. In JavaScript this is done with the `concat` method rather than an operator.
- All elements in an Elm list must have a consistent type. For example, it can be a "list of numbers" or a "list of strings," but not a "list where strings and numbers intermingle." (Making a list containing both strings and numbers involves first creating wrapper elements for them, using a feature called *union types* that we'll cover in Chapter 3.)

Although Elm supports both (immutable) lists and (also immutable) arrays, lists are used far more often because they have better performance characteristics in typical Elm use cases. We'll get more into the performance differences between linked lists and arrays in Chapter 10.

Here are some examples of how Elm lists and JavaScript arrays differ.

**Table 1.4 Contrasting JavaScript Arrays and Elm Lists**

JavaScript Array	Elm List
<code>[ 1, 2, 3 ].length</code>	<code>List.length [ 1, 2, 3 ]</code>
<code>[ "one fish", "two fish" ][0]</code>	<code>List.head [ "one fish", "two fish" ]</code>
<code>[ "one fish", "two fish" ][1]</code>	No arbitrary position-based element access
<code>[ 1, 2 ].concat([ 3, 4 ])</code>	<code>[ 1, 2 ] ++ [ 3, 4 ]</code>
<code>[ 1, 2 ].push(3)</code>	Cannot be modified; use e.g. <code>append</code> instead
<code>[ 1, "Paper", 3 ]</code>	All elements in a list must have a consistent type

Let's focus on that last one. Why must all elements in an Elm list have a consistent type?

To understand how this requirement benefits us, let's consider the `List.filter` function, which works like the `String.filter` function we used earlier.

Recall that `String.filter` takes a function which returns `True` when the given character should be kept, and `False` when it should be dropped. `List.filter` differs only in that the function you provide doesn't necessarily receive characters—instead it receives elements from the list, whatever they may be.

Let's see that in action. Quick! To `elm-repl`!

**TIP** Whereas before we had to `import String`, the `List` module is one of a select few modules that gets imported by default. We don't need to write `import List` to reference it.

### Listing 1.14 Filtering lists

```
> List.filter (\char -> char /= '-') [ 'Z', '-', 'Z' ] ❶
['Z','Z']

> List.filter (\str -> str /= "-") [ "ZZ", "-", "Top" ] ❷
["ZZ","Top"]

> import Char
> List.filter Char.isDigit [ '7', '-', '9' ] ❸
['7','9']

> List.filter (\num -> num % 2 == 1) [ 1, 2, 3, 4, 5 ] ❹
[1,3,5]
```

- ❶ Same function we passed to `String.filter` earlier
- ❷ Strings instead of characters
- ❸ Works just like with `String.filter`
- ❹ Keep only the odd numbers

Here's how we would rewrite that last line of code in JavaScript:

```
[ 1, 2, 3, 4, 5 ].filter(function(num) { return num % 2 === 1; })
```

This looks straightforward enough, but JavaScript arrays permit inconsistent element types. Without looking it up, can you guess what happens if we change it to the following?

```
[ 1, "2", "cat", 4, "5", "" ].filter(function(num) { return num % 2 === 1; })
```

Will it crash? Will it happily return numbers? What about strings? It's a bit of a head-scratcher.

Because Elm requires consistent element types, this is a no-brainer: in Elm it would be an error. Even better, it would be an error at build time—meaning you can rest easy knowing whatever surprises would result from executing this code will not inflict pain on your users. Requiring consistent element types means all lists in Elm guarantee this level of predictability.

By the way, the above `filter()` call successfully returns `[ 1, "5" ]`. (Like, *duh*, right?)

## 1.4.2 Records

We've now seen how JavaScript's mutable arrays resemble Elm's immutable lists. In a similar vein, JavaScript's mutable objects resemble Elm's immutable *records*.

**DEFINITION** A *record* is a collection of named fields, each with an associated value.

Whereas array and list literals between the two languages are syntactically identical, where JavaScript object literals use `:` to separate fields and values, Elm record literals use `=` instead.

Let's get a taste for some of their other differences.

JavaScript Object	Elm Record
<code>{ name: "Li", cats: 2 }</code>	<code>{ name = "Li", cats = 2 }</code>
<code>({ name: "Li", cats: 2 }).toString()</code>	<code>toString { name = "Li", cats = 2 }</code>
<code>({ name: "Li", cats: 2 }).cats</code>	<code>(&lt;{ name = "Li", cats = 2 }).cats</code>
<code>({ name: "Li", cats: 2 })["cats"]</code>	Fields can only be accessed directly, using a dot
<code>({ name: "Li", cats: 2 }).cats = 3</code>	Cannot be modified. (New cat? New record!)
<code>{ NAME: "Li", CATS: 2 }</code>	Fields have the same naming rules as constants
<code>({ name: "Li", cats: 2 }).__proto__</code>	No prepackaged fields, only the ones you define
<code>Object.keys({ name: "Li", cats: 5 })</code>	No listing of field names is available on demand
<code>Object.prototype</code>	Records have no concept of inheritance

Wow—compared to objects, records sure don't do much! It's like all they do is sit around holding onto the data we gave them. (Yep.) Personally I've found Elm's records a welcome reprieve from the intricacies of JavaScript's `this` keyword.

## RECORD UPDATES

*Record updates* let us concisely obtain a new record by copying the old one and changing only the specified values. (As we will see in Chapter 10, behind the scenes Elm does not actually copy the entire record—that would be slow!—but rather only the parts that will be different.)

Let's use this technique to represent someone obtaining an extra cat, going from `{ name = "Li", cats = 2 }` to `{ name = "Li", cats = 3 }` by way of a record update.

### Listing 1.15 Record updates

```
> catLover = { name = "Li", cats = 2 }
{ name = "Li", cats = 2 }

> catLover
{ name = "Li", cats = 2 }

> withThirdCat = { catLover | cats = 3 } ❶
{ name = "Li", cats = 3 }

> withThirdCat
{ name = "Li", cats = 3 }

> catLover ❷
{ name = "Li", cats = 2 } ❷

> { catLover | cats = 88, name = "LORD OF CATS" } ❸
{ name = "LORD OF CATS", cats = 88 }
```

- ❶ Record update syntax
- ❷ Original record unmodified!

### 3 Update multiple fields (order doesn't matter)

Record updates let us represent this incremental evolution without mutating our records or recreating them from scratch. In Chapter 2 we'll represent our application state with a record, and use record updates to make changes based on user interaction.

## 1.4.3 Tuples

Lists let us represent collections of **varying size**, whose elements share a **consistent type**. Records let us represent collections of **fixed fields**, but whose corresponding values may have **varied types**.

*Tuples* introduce no new capabilities to this mix, as there is nothing a tuple can do that a record couldn't. Compared to records, though, what tuples bring to the party is conciseness.

**DEFINITION** A *tuple* is a record-like value whose fields are accessed by position rather than by name.

In other words, tuples are for when you want a record, but don't want to bother naming its fields. They are often used for things like key-value pairs where writing out `{ key = "foo", value = "bar" }` would add verbosity but not much clarity.

Let's try some out!

### Listing 1.17 Using Tuples

```
> ( "Tech", 9 )
("Tech",9)

> fst ( "Tech", 9 ) ❶
"Tech"

> snd ( "Tech", 9 ) ❷
9
```

- ❶ Return first element (only works on 2-element tuples)
- ❷ Return second element (only works on 2-element tuples)

You can only use the `fst` and `snd` functions on tuples that contain two elements. If they have more than two, you can use *tuple destructuring* to extract their values.

**DEFINITION** *Tuple destructuring* extracts the values inside a tuple and assigns them to constants in the current scope.

Let's use tuple destructuring to implement a function that takes a tuple of three elements.

**Listing 1.18 Tuple Destructuring**

```

> multiply3d ( x, y, z ) = x * y * z ❶
<function>

> multiply3d ( 6, 7, 2 )
84

> multiply2d someTuple = let ( x, y ) = someTuple in x * y ❷
<function>

```

- ❶ Destructuring a tuple into three constants: x, y, and z
- ❷ Destructuring a tuple inside a let-expression

As demonstrated in Listing 1.18, once you have named the values inside the tuple, you can use them just like you would any other constant.

**TIP** Mind the difference between a tuple and a parenthetical function call! `( foo, bar )` is a tuple, whereas `( foo bar )` is a call to the `foo` function passing `bar` as an argument. A simple mnemonic to remember the difference is “comma means tuple.”

**Table 1.5 Comparing Lists, Records, and Tuples**

List	Record	Tuple
Variable Length	Fixed Length	Fixed Length
Can Iterate Over	Cannot Iterate Over	Cannot Iterate Over
No Names	Named Fields	No Names
Immutable	Immutable	Immutable

Since any tuples can be represented (more verbosely) using a record instead, it’s often better to refactor long tuples—say, with more than three elements—into records. Choose tuples or records based on whichever would yield more readable code; their performance characteristics are equivalent.

## 1.5 Summary

We’re off to a fantastic start! First we discussed some of the toughest problems Web programmers face: crashing is too easy in JavaScript, and maintenance is too error-prone. Then we learned how Elm addresses these problems, with a design that prioritizes maintainability and a helpful compiler that catches would-be runtime exceptions before they can cause user pain. From there we dove in and wrote our first Elm code in `elm-repl`.

Here is a brief review of things we covered along the way.

- The `++` operator combines strings and lists, whereas the `+` operator is for addition only.
- Double quotes refer to strings. Single quotes refer to individual UTF-8 characters.

- *let-expressions* introduce scoped constants to an expression.
- There is no concept of “truthiness” in Elm, just `True` and `False`.
- `if foo /= bar then "different" else "same"` is an *if-expression*. Like JavaScript ternaries, *if-expressions* require an `else` branch and always evaluate to a value.
- Lists like `[ 3, 1, 4 ]` are immutable. Their elements must share a consistent type.
- `List.filter (\num -> num < 0) numbersList` returns a list containing all the negative numbers in the original `numbersList`
- `catLover = { name = "Li", cats = 2 }` assigns a record to the constant `catLover`. Once assigned, constants cannot be reassigned.
- `{ catLover | cats = 3 }` returns a new record that is the same as the `catLover` record, except the `cats` value is now 3.
- `( foo, bar )` deconstructs a tuple such as `( 2, 3 )`. In this example, `foo` would be 2 and `bar` would be 3.

Table 1.6 summarizes some of the differences between JavaScript and Elm.

**Table 1.6 Differences between JavaScript and Elm**

JavaScript	Elm
<code>// This is an inline comment</code>	<code>-- This is an inline comment</code>
<code>/* This is a block comment */</code>	<code>{- This is a block comment -}</code>
<code>true &amp;&amp; false</code>	<code>True &amp;&amp; False</code>
<code>"Ahoy, " + "World!"</code>	<code>"Ahoy, " ++ "World!"</code>
<code>"A spade" === "A spade"</code>	<code>"A spade" == "A spade"</code>
<code>"Calvin" !== "Hobbes"</code>	<code>"Calvin" /= "Hobbes"</code>
<code>Math.pow(2, 11)</code>	<code>2 ^ 11</code>
<code>Math.trunc(-49 / 10)</code>	<code>-49 // 10</code>
<code>n % 2 === 1 ? "odd" : "even"</code>	<code>if n % 2 == 1 then "odd" else "even"</code>
<code>nums.filter(function(n) { ... })</code>	<code>List.filter (\n -&gt; n % 2 == 1) nums</code>
<code>function pluralize(s, p, c) { ... }</code>	<code>pluralize singular plural count = ...</code>

We also learned about several differences between normal functions and operators:

**Table 1.7 Differences between normal functions and operators**

Function	How to identify one	Calling style	Examples
Normal	Name begins with a letter	<i>prefix-style</i>	<code>negate</code> , <code>not</code> , <code>toString</code>
Operator	Name has no letters or numbers	<i>infix-style</i>	<code>(++)</code> , <code>(*)</code> , <code>(==)</code>

In Chapter 2 we'll expand on what we've learned here to create a working Elm application.

Let's go build something!